




Discrete-Event Modeling

Discrete-event simulation

- We want to model a system using a **discrete-event simulator**
 - Events are instantaneous time occurrences at which the system *state* changes





Discrete-event simulation


- We want to model a system using a **discrete-event simulator**
 - Events are instantaneous time occurrences at which the system *state* changes
 - We are interested in representing the system in these instants of time
 - **Time advances through discrete steps from one event to the next**
 - The way the system evolves depends on the nature of the **intervals between events**, whose duration can be either **deterministic or stochastic**



Discrete-event simulation

- Since we are interested in the system just in the moments at which its state changes, there is no need to model the system between state changes
- We use a **next event** approach: the system time jumps forward with discrete steps to the time instant at which the next state change (event) happens

Example



- The events occur in time instants t_1, t_2, \dots
- Time advances corresponding to these instants



Discrete-event simulation

- Definitions (reminder)
 - *Entities*: objects of interest in the system
 - *Event*: an instantaneous occurrence that changes the system state
 - *Activity*: a time period dedicated to a specific operation
 - *Process*: a sequence of events, ordered in time, involving a single entity



Discrete-event simulation

- We have some alternative solutions for discrete-event modeling
 - we can focus on the actions to be performed when each event happens -> *event scheduling*
 - we consider all the actions of each entity while it is inside the system, focusing on its relations with other entities -> *process interaction*



Event scheduling

- *To schedule*
 - To plan the future execution of an operation
- The *event scheduling* approach requires that we decide when in the future an event must happen and we schedule its execution
- It is the most common approach
 - Supported by almost any simulation language
 - Easy coding



Event scheduling

- We need to identify or define
 - the **entities** in the system
 - the **results** we want to obtain from the simulation
 - the **state variables**
 - the **events**, i.e., what makes the state variables change



Event scheduling

- For each event, **we must define the sequence of operations** to be executed when the event happens
- These operations depend on the state of
 - **the entity that triggered the event**
 - **other entities in the system**



Event scheduling

- A special and very important operation executed when an event is triggered is *the scheduling of new events*
- In this way, starting from **few events scheduled at the beginning of the simulation**, the system will naturally evolve through the cause/effect relations among events




Event scheduling

- Let us use a concrete model: **the queue model**
- The best way to gain a deep understanding of the **next-event scheduling approach** is to build step-by-step the simulation model for a key element in telecommunication networks: **the queue**

Queuing systems

- Systems in which
 - **Customers** arrive to receive **service**
 - One or more **servers** provide service





Model of a queuing system

- We consider a queuing system
- In such a system users (**customers or clients**) **receive in an ordered way** a service, provided by one or more agents (**servers**) according to a specific service discipline
- It is the most common (sub)system we can find studying a telecommunication network



Model of a queuing system

- Which are the elements/processes?
 - the clients requiring service
 - the servers providing it
- What are we interested in?
 - average time to be served (waiting time)
 - average number of clients waiting for service
 - when there is limited space for the waiting clients, the probability that a client might not be served



Model of a queuing system

- Which are the state variables?
 - the number of clients in the system
- Which are the events?
 - the arrival of a client
 - the end of the service of a client



Model of a queuing system

- Let us start with a very simple version of our model, improving it later with incremental steps after identifying any detail needed and missing

- Version 1.0

- Arrival

- If the server is idle, occupy it
- Otherwise, enter the waiting line

- End of Service

- If there is someone in the waiting line, serve it
- otherwise, become idle



Model of a queuing system

- Version 1.0 is very simplistic
- The relations between events are missing
 - **End of Service** events are caused by the presence of clients in the system and by their services
- We need to explicit these relations



Model of a queuing system

- Version 2.0

- Arrival

- If the server is idle
 - Make the server busy
 - Schedule the end of service
- else
 - Wait in line



Model of a queuing system

- Version 2.0
- End of Service
 - If someone is in the waiting line
 - Select next in line
 - Schedule the end of service
 - else
 - Make the server idle



Model of a queuing system

- We have pointed out how the end of each service is planned at its start
- We are still missing a few key details in our simulation model
 - Each service has its own **duration** that must be chosen according to the requirements of the considered system
 - We may need to store information about the waiting clients in suitable data structures



Model of a queuing system

- Version 3.0

- Arrival

- If the server is idle

- Make the server busy
 - Determine the service time T_s
 - Schedule the end of service at time $T_{curr} + T_s$

- else

- Create a record for the client
 - Insert the record in the queue



Model of a queuing system

- Version 3.0

- End of Service

- If someone is in the waiting line
 - Select a client to be served
 - Extract its record from the queue
 - Determine service time T_s
 - Schedule the end of service at time $T_{curr} + T_s$
- else
 - Make the server idle



Model of a queuing system

- We are still missing a key detail

Who is controlling the arrival of clients?

- Clients arrival is usually characterized as a random process, e.g. through the distribution of the inter-arrival time
- The most common solution to manage the arrival of clients in the system is by self-sustaining
 - At each arrival, the next one is scheduled
- We also need to develop a correct and coherent management of the records for the clients



Model of a queuing system v4.0

■ Arrival

- Compute the inter-arrival time T_{ia} for next client
- Schedule an arrival at time $T_{curr} + T_{ia}$
- Create a record for the client
- Insert the record in the queue
- If the server is idle
 - Make the server busy
 - Determine the service time T_s
 - Schedule the end of service at time $T_{curr} + T_s$



Model of a queuing system v4.0

■ End of Service

- Extract from the queue the record for the client who has just been served
- Destroy the record
- If someone is in the waiting line
 - Select a client to be served
 - Determine its service time T_s
 - Schedule the end of service at time $T_{curr} + T_s$
- else
 - Make the server idle



Model of a queuing system

- We can further increase the details of the model, considering
 - the queuing policy
 - detailed management of the record for the clients
 - the operations needed to measure the desired performance indexes
- Measurements are quite a critical issue
 - Where and how do we carry them out?



Process interaction

- This modeling approach is focused on the actions performed by each entity during its lifetime in the system
- A *process* is a sequence of events and activities involving a single entity
- The behavior of each entity depends on the actions and the state of the other entities in the system
- In the model we need to highlight the **interaction** between entities and, as a consequence, between the processes they are represented from



Process interaction

- The process represents the actions of the entity during its lifetime
- The events are now representing **suspensions** and **reactivations** in the execution flow of the entity
- Two classes of events:
 - *unconditional* -> the time at which the event happens is known in advance
 - *conditional* -> the time at which the event actually happens depends on external conditions (e.g. the state of other entities)



Process interaction

- The server model is a non-terminating loop: we need just to schedule its execution once, when the simulation starts
- Several client processes may be active at the same time, one for each waiting client, but there is always a single server process
- Suspension/reactivation points are indicated by the blue text: they are the moments at which the control flow moves from one process to the other



Process interaction

- Process interaction is a very powerful modeling approach
 - The behavioral description of each entity is concentrated in a single block, instead of being distributed among several events
- The compact representation of the model hides the implementation complexity due to
 - conditional/unconditional events
 - reactivation points of the execution flow



Wrap-up

- In discrete-event simulation, events are instants in which the system state changes
- Next event approach: times evolves according to the occurrence of events
- Two main approaches:
 - Event scheduling:
 - model actions performed at each event
 - Process interaction
 - model actions of each entity